




**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Multi-Node Distributed Training with tf.keras

CSCS TensorFlow 2020

Henrique Mendonça and Rafael Sarmiento, CSCS

7-8 Sep 2020



# Motivation: Why distribute?

- Why should we use multiple GPU's to train a DL model?



# Motivation: Why distribute?

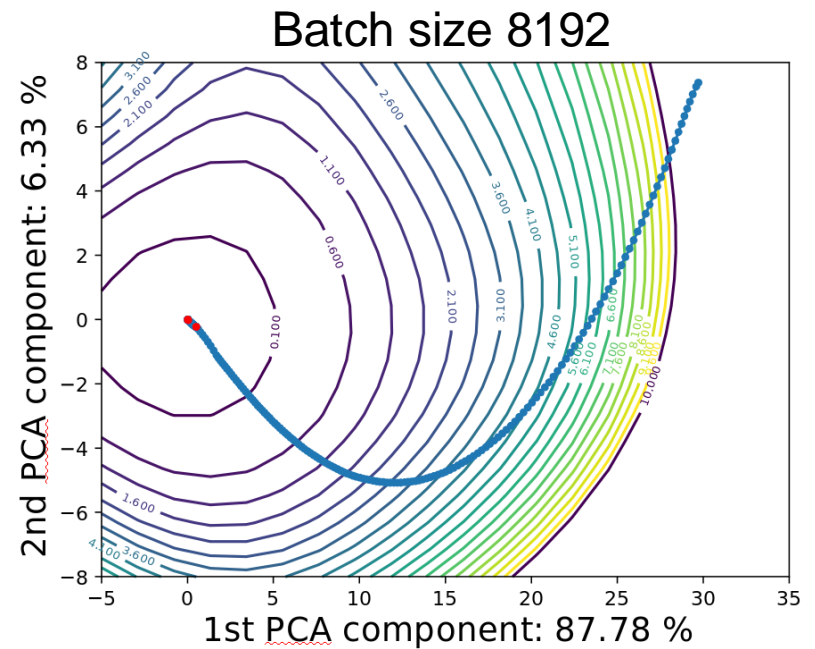
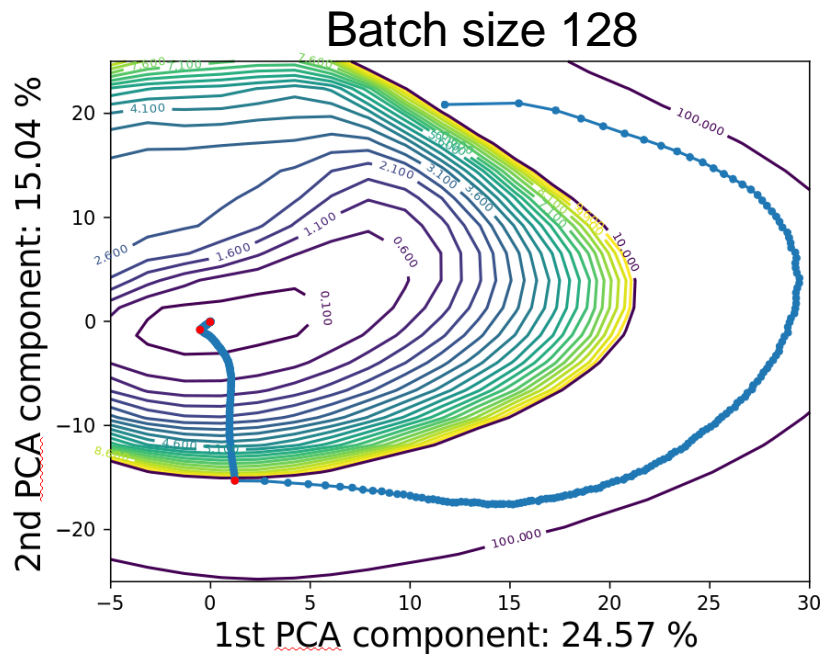

- Because it's faster
  - Allow quick experimentation of new ideas and models



- Allows training models larger than memory
  - Out-of-Core learning: Not covered in this course
- May allow better accuracy, especially in larger models. Why?


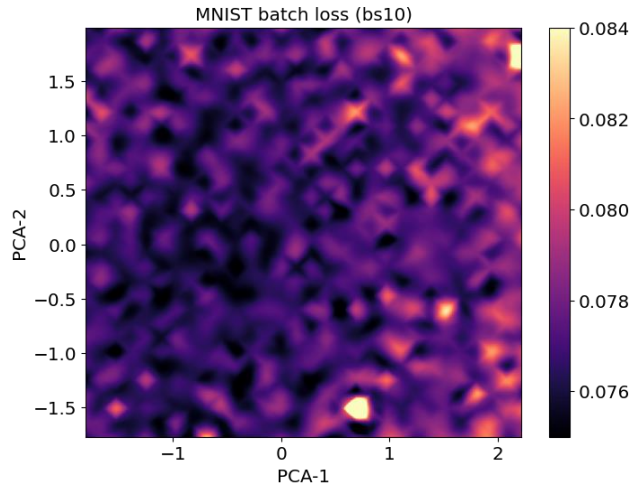
# Motivation: Sample noise

Visualizing the Loss Landscape of Neural Nets  
<https://arxiv.org/pdf/1712.09913.pdf>


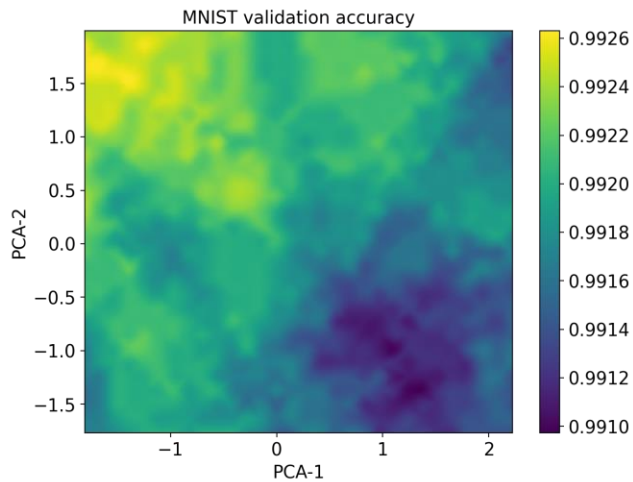


# Motivation: Sample noise

- Training:



- Validation:



# Distributed training with TF Keras 2.3

- TensorFlow introduces tf.distribute.Strategy
  - Wraps model in a distributed scope
- Multi-GPU -> tf.distribute.MirroredStrategy
  - Single node/worker

```
strategy = tf.distribute.MirroredStrategy()  
with strategy.scope():  
    model = build_and_compile_model()  
model.fit(dataset, epochs, steps_per_epoch)
```

- NCCL AllReduce by default
- Automatic data sharding across GPU's

# Multi-worker Distribution

- Creates some additional complexity
  - external network communication
  - separated OS
  - separated processes
    - Facilitated by `ipyparallel` magic on JupyterLab





# Multi-worker Distribution

- [tf.distribute.experimental.MultiWorkerMirroredStrategy](#)
- Communication:
  - NCCL AllReduce for all-reduce (if available)
  - Ring algorithm for all-gather
  - Includes fault tolerance when using [BackupAndRestore](#)
  - Datasets are automatically sharded across nodes
- Cluster Resolver:
  - defaults to TFConfig

```
os.environ['TF_CONFIG'] = '{  
    "cluster": {"worker": ["nid01111:8888", "nid02222:8888"]},  
    "task": {"type": "worker", "index": "0"}  
}'
```



# Multi-worker distribution with SLURM

- TensorFlow 2.2+

```
tf.distribute.cluster_resolver.SlurmClusterResolver(  
    port_base=8888, auto_set_gpu=True, rpc_layer='grpc',  
    jobs=None, gpus_per_node=None, gpus_per_task=None,  
    tasks_per_node=None  
)
```

- All parameters are automatically queried from SLURM



# Multi-worker distribution with SLURM

```
%%px
strategy = tfd.experimental.MultiWorkerMirroredStrategy(
    cluster_resolver=tfd.cluster_resolver.SlurmClusterResolver(),
    communication=tfd.experimental.CollectiveCommunication.NCCL,
)
with strategy.scope():
    model = build_and_compile_model()
model.fit(dataset, epochs=N, steps_per_epoch=M)
```

# Done!



# Multi-worker distribution with SLURM

- Sharding
  - Both Training and Validating datasets are automatically distributed across nodes
    - File sharding is used when possible (TFrecords)
- Fault Tolerance
  - The BackupAndRestore callback saves a checkpoint every epoch and restore when needed
- `model.predict()` is not supported
  - Use a single node instead

# Multi-worker distribution with SLURM

- Practise
  - Run MNIST training and inference on your GPU's
    - [03\\_mnist\\_tf.distribute.ipynb](#)

# Batch Norm Synchronisation

- BN: Exponential Moving Average per channel

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

$$\hat{x}_i^{(k)} = \frac{x_i^{(k)} - \mu_B^{(k)}}{\sqrt{\sigma_B^{(k)^2} + \epsilon}}$$


$$y_i^{(k)} = \gamma^{(k)} \hat{x}_i^{(k)} + \beta^{(k)}$$

- [tf.keras.layers.experimental.SyncBatchNormalization](#)
  - AllReduce across BN layers during forward pass
  - Synchronizes all statistics before autodiff


# ImageNet

- Transfer Learning made Deep Learning accessible

ImageNet-2012: 1,281,167 samples



1st convolutional layer of  
SeResNeXT101.32x4d (ImageNet)

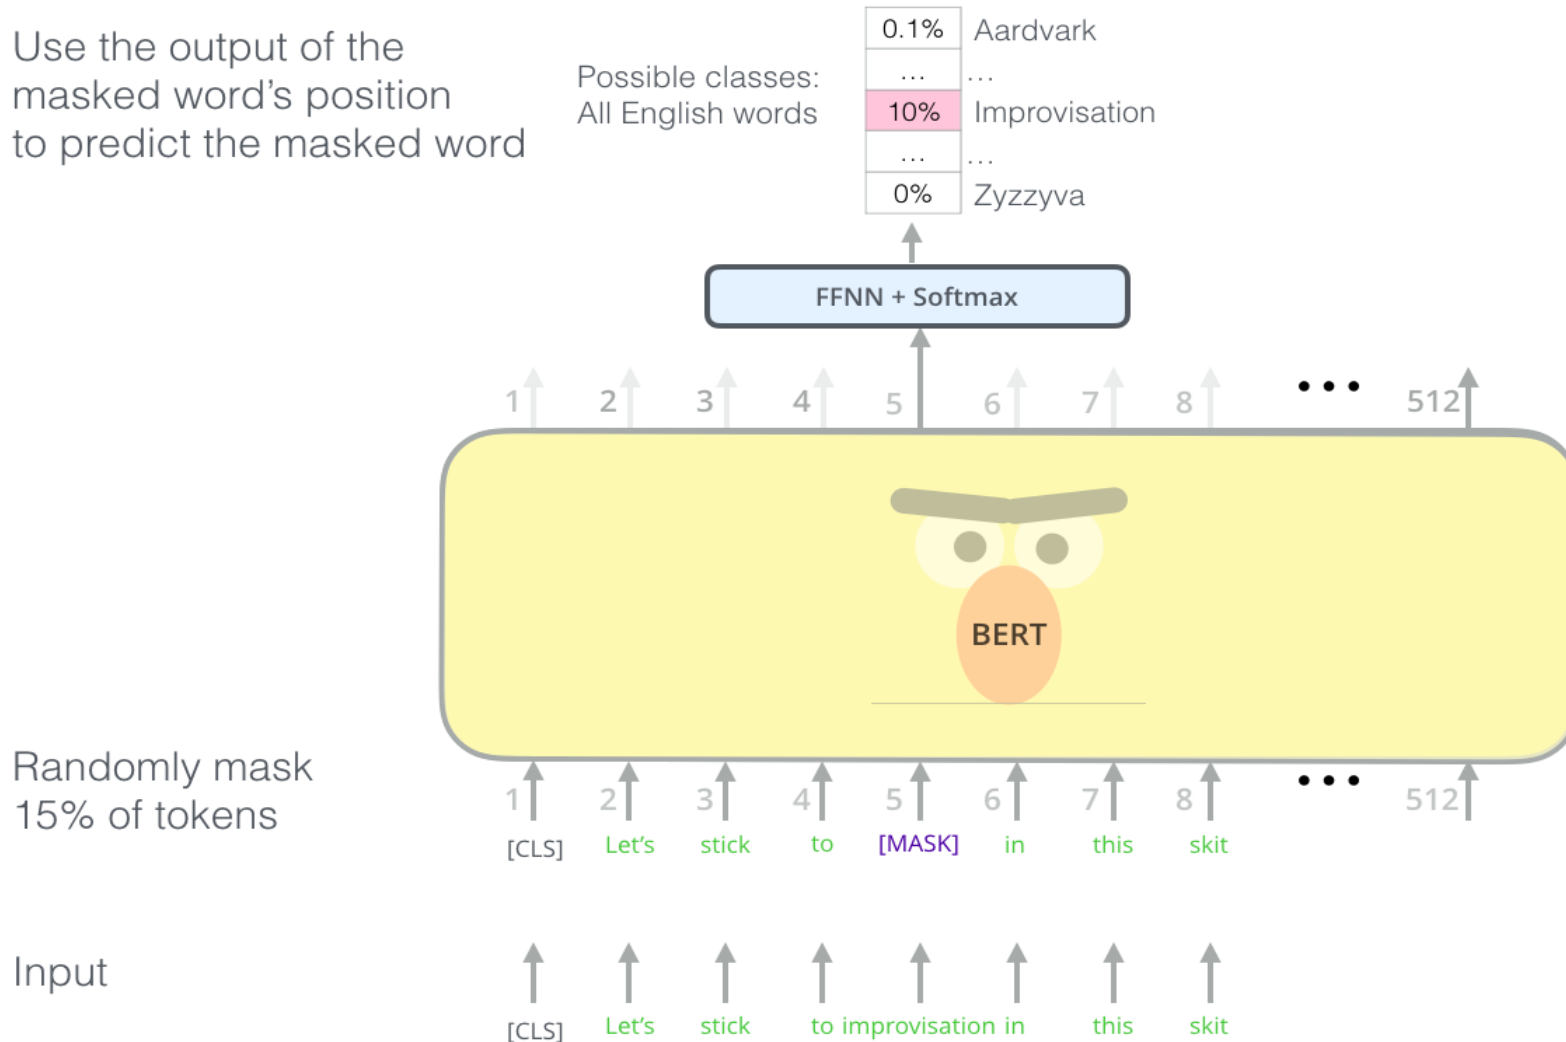


# BERT: LM Encoder-Decoder

Use the output of the masked word's position to predict the masked word

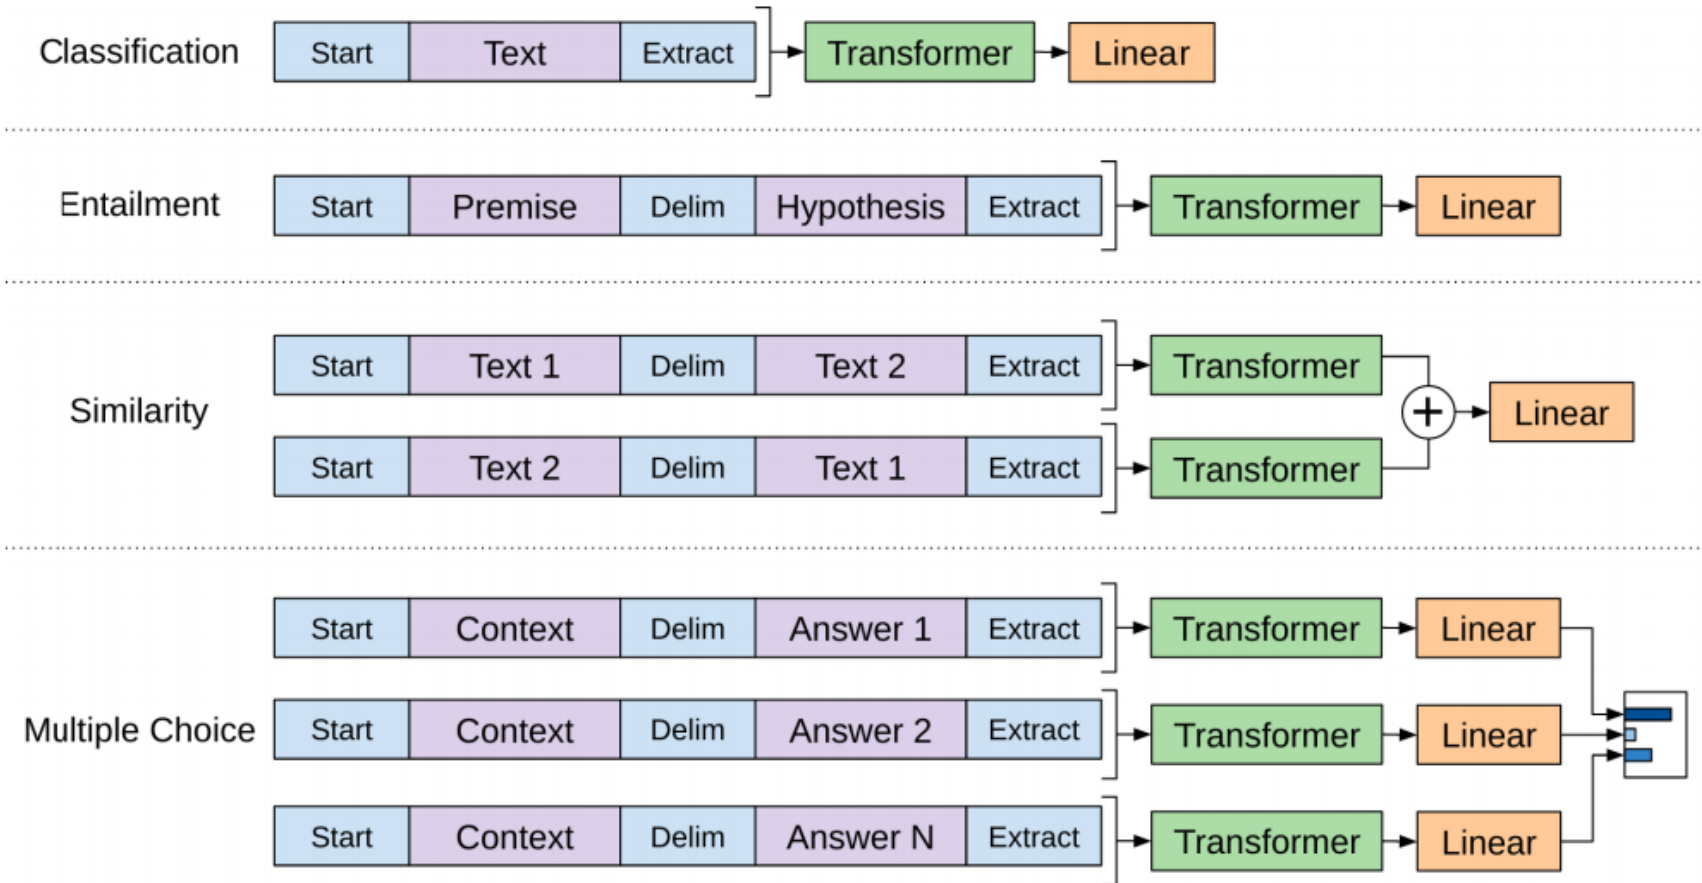
Possible classes:  
All English words

|      |               |
|------|---------------|
| 0.1% | Aardvark      |
| ...  | ...           |
| 10%  | Improvisation |
| ...  | ...           |
| 0%   | Zyzzzyva      |





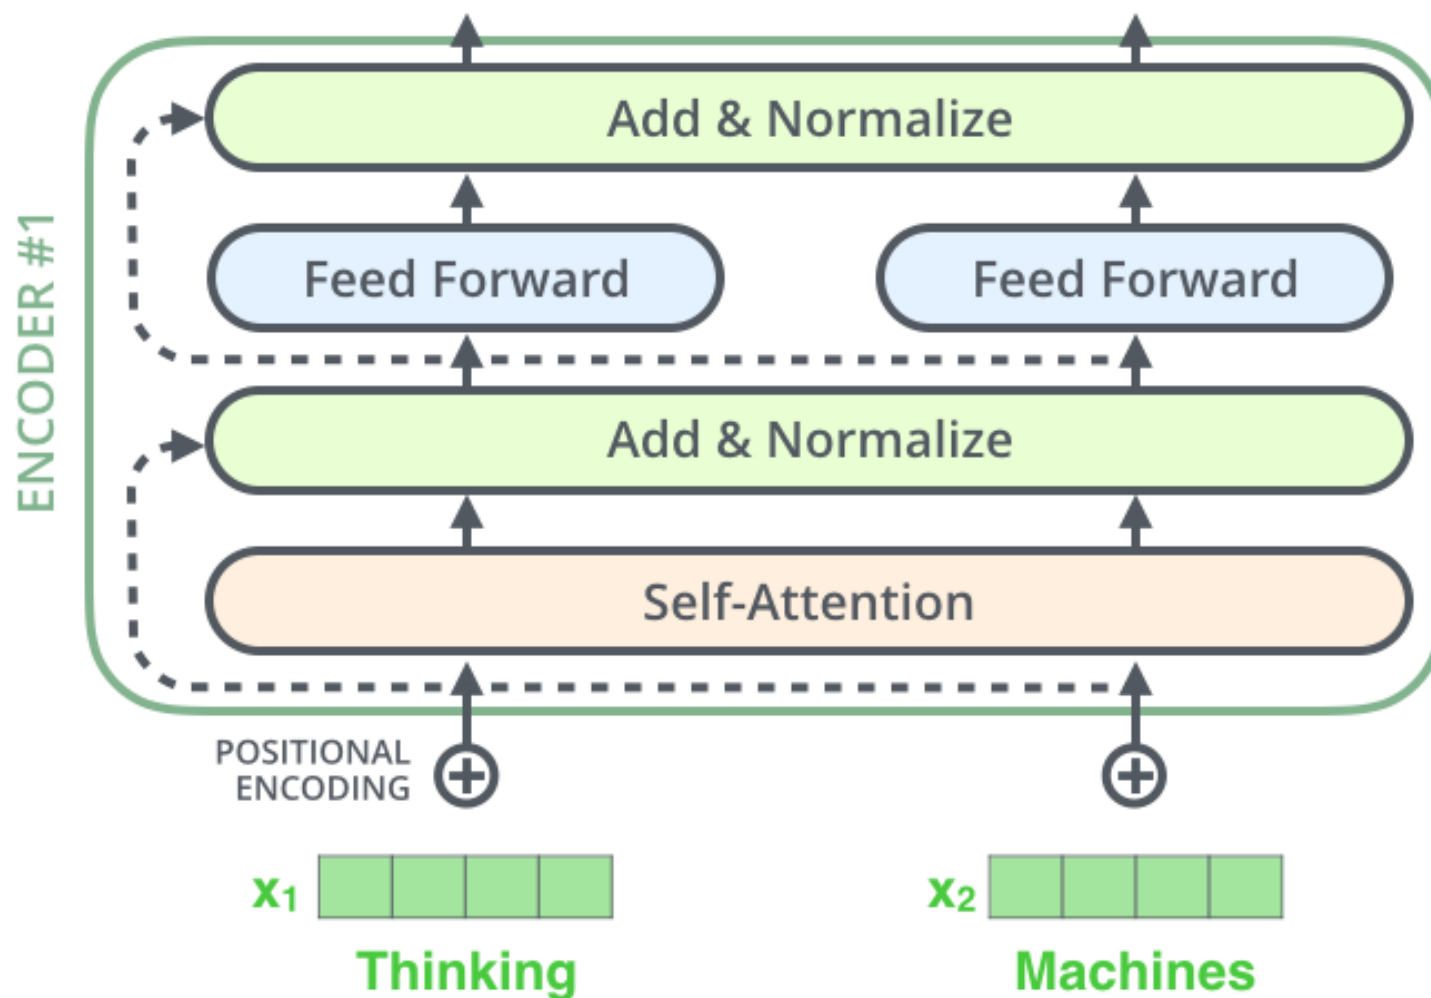
# BERT: Encoder Fine-Tuning



Source <http://jalammar.github.io/illustrated-bert/>

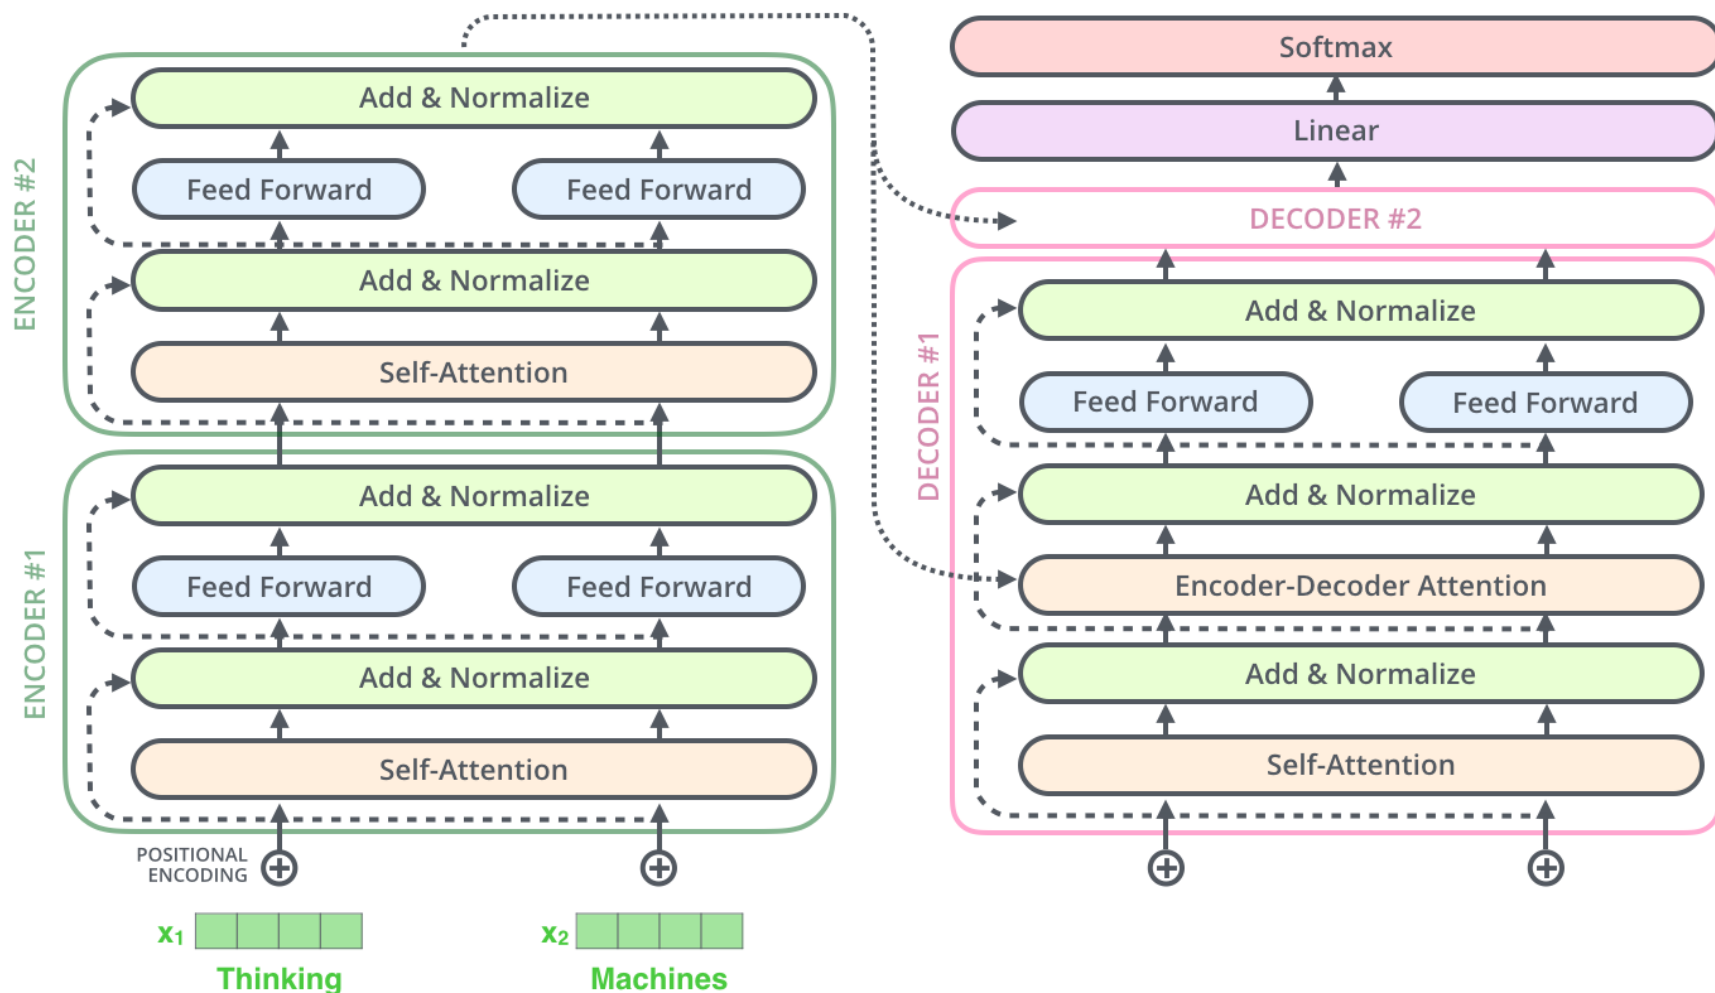
BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding <https://arxiv.org/abs/1810.04805>

# Transformer: Encoder Blocks



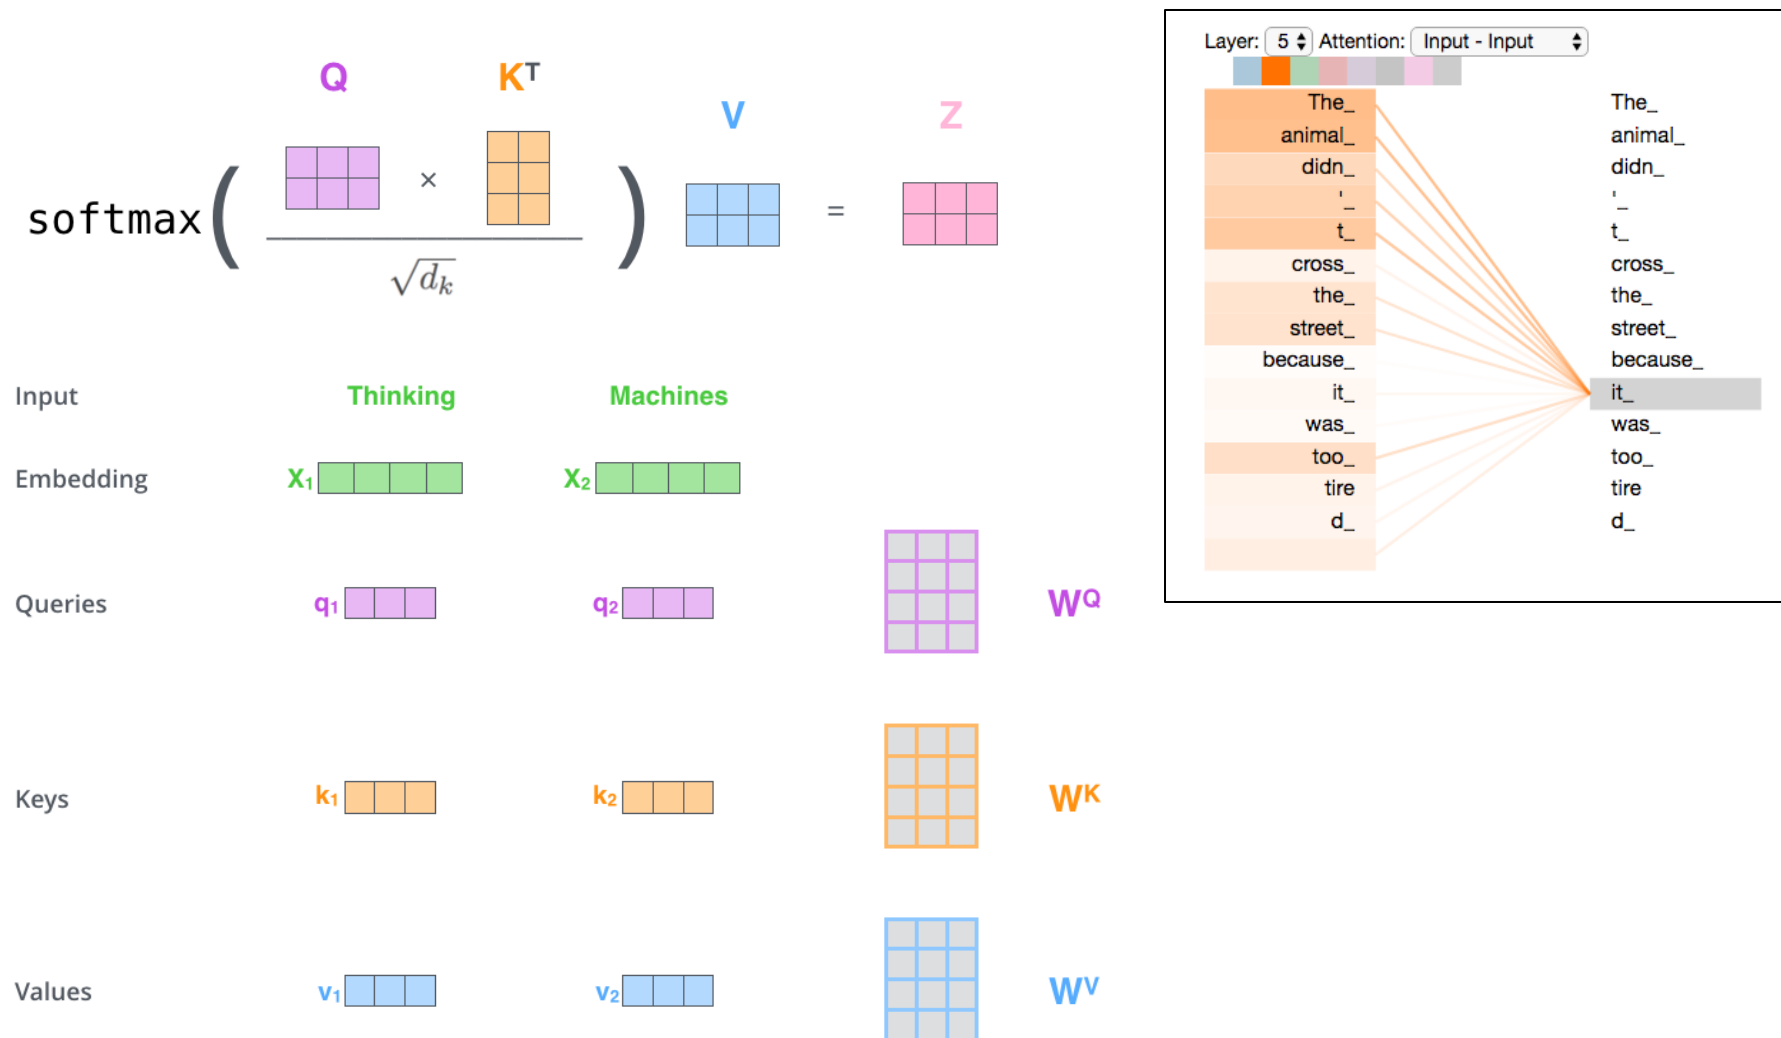
Source <https://jalammar.github.io/illustrated-transformer>  
Attention Is All You Need <https://arxiv.org/pdf/1706.03762.pdf>

# Transformer: LM Encoder-Decoder




Source <https://jalammar.github.io/illustrated-transformer>  
Attention Is All You Need <https://arxiv.org/pdf/1706.03762.pdf>

# Transformer: Self-Attention at a High Level



Source <https://jalammar.github.io/illustrated-transformer>  
 Attention Is All You Need <https://arxiv.org/pdf/1706.03762.pdf>



**Thank you for your attention.**